

18.s096 pset 7 - Dimitris Koutentakis

May 8, 2018

1 Problem 1

1.1 (a)

```
In [2]: library("mcsn")
        data(challenger)
        fit.logistic<-glm(oring ~ temp, data=challenger, family=binomial(link="logit"))
        fit.logistic
        summary(fit.logistic)
```

```
Call: glm(formula = oring ~ temp, family = binomial(link = "logit"),
  data = challenger)
```

Coefficients:

(Intercept)	temp
15.0429	-0.2322

Degrees of Freedom: 22 Total (i.e. Null); 21 Residual

Null Deviance: 28.27

Residual Deviance: 20.32 AIC: 24.32

Call:

```
glm(formula = oring ~ temp, family = binomial(link = "logit"),
  data = challenger)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.0611	-0.7613	-0.3783	0.4524	2.2175

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	15.0429	7.3786	2.039	0.0415 *
temp	-0.2322	0.1082	-2.145	0.0320 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 28.267 on 22 degrees of freedom
Residual deviance: 20.315 on 21 degrees of freedom
AIC: 24.315

Number of Fisher Scoring iterations: 5

From the above, we get the MLE values of α and β . In specific, we have:

$$\hat{\alpha}_{MLE} = 15.0429, \hat{\beta}_{MLE} = -0.2322$$

The we have that the 95% confidence intervals are:

$$I_{\alpha} = [\hat{\alpha}_{MLE} - 1.96 \cdot 7.3786, \hat{\alpha}_{MLE} + 1.96 \cdot 7.3786] \quad (1)$$

$$= [0.58, 29.505] \quad (2)$$

$$I_{\beta} = [\hat{\beta}_{MLE} - 1.96 \cdot 0.1082, \hat{\beta}_{MLE} + 1.96 \cdot 0.1082] \quad (3)$$

$$= [-0.444, -0.02] \quad (4)$$

1.2 (b)

```
In [3]: set.seed(1)
        Nsim=10^4
        x=challenger$temp
        y=challenger$oring
        sigmaa=5 ; sigmab=5/sd(x)
        lpost=function(a,b){sum(y*(a+b*x)-log(1+exp(a+b*x)))+
                                dnorm(a,sd=sigmaa,log=TRUE)+dnorm(b,sd=sigmab,log=TRUE)}
        # Initialize a and b to equal the MLEs
        beta=as.vector(fit.logistic$coefficients)
        a=b=rep(0,Nsim)
        a[1]=beta[1]
        b[1]=beta[2]
        #As scale for the proposal densities consider the square root of the
        # cov.unscaled from the ml fit to the logistic model
        fit.logistic.summary<-summary(fit.logistic)
        scala=sqrt(fit.logistic.summary$cov.unscaled[1,1])
        scalb=sqrt(fit.logistic.summary$cov.unscaled[2,2])
        for (t in 2:Nsim){
            propa=a[t-1]+sample(c(-1,1),1)*rexp(1)*scala
            if (log(runif(1))<lpost(propa,b[t-1])- lpost(a[t-1],b[t-1]))
                a[t]=propa else
```

```

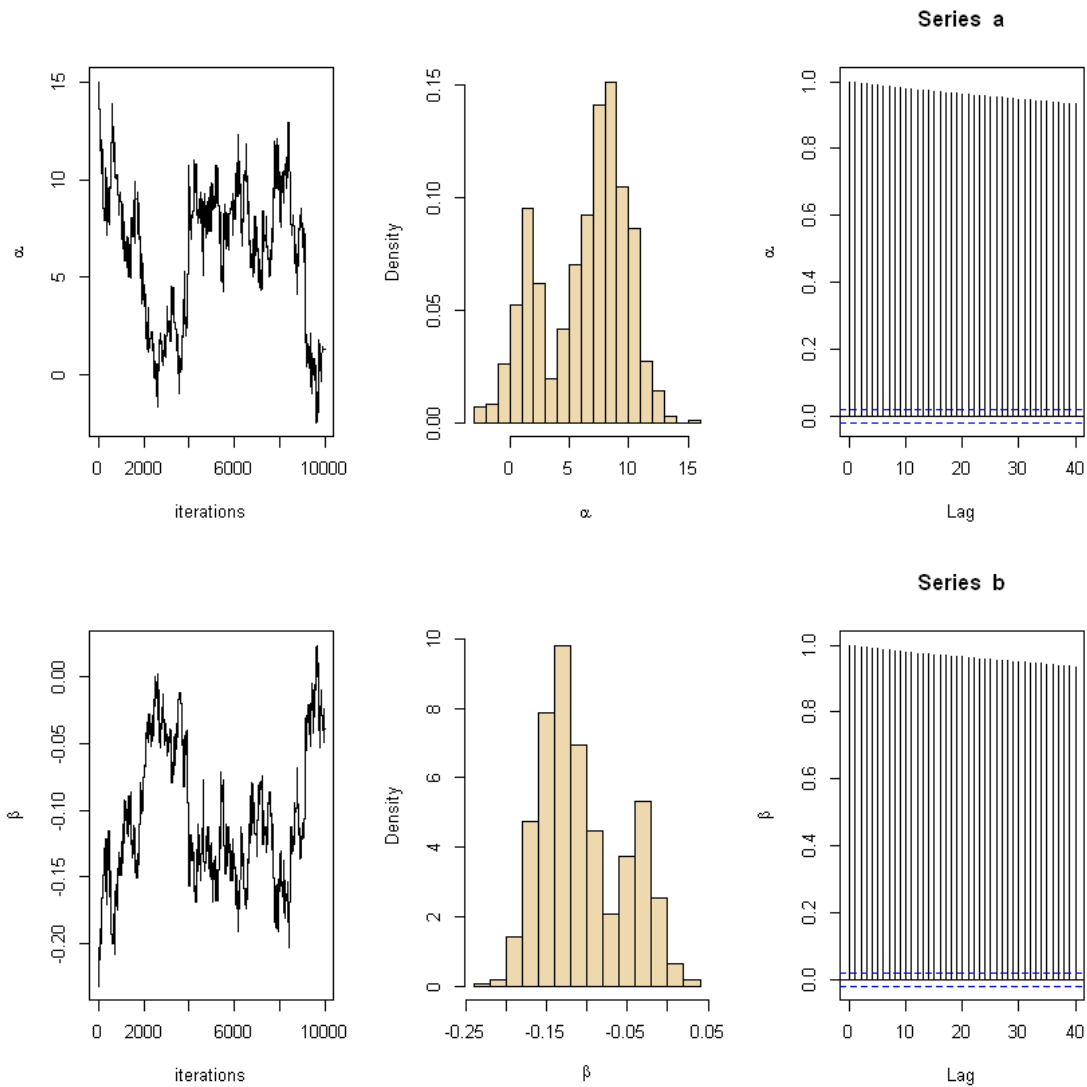
        a[t]=a[t-1]
        propb=b[t-1]+sample(c(-1,1),1)*rexp(1)*scalb
        if (log(runif(1))<lpost(a[t],propb)- lpost(a[t],b[t-1]))
            b[t]=propb
        else b[t]=b[t-1]
    }

In [4]: print(length(unique(a))/Nsim)
        print(length(unique(b))/Nsim)

[1] 0.1001
[1] 0.0987

In [5]: par(mfrow=c(2,3))
        plot(a,type="l",xlab="iterations",ylab=expression(alpha))
        hist(a,prob=TRUE,col="wheat2",xlab=expression(alpha),main="")
        acf(a,ylab=expression(alpha))
        plot(b,type="l",xlab="iterations",ylab=expression(beta))
        hist(b,prob=TRUE,col="wheat2",xlab=expression(beta),main="")
        acf(b,ylab=expression(beta))

```




```
In [6]: print (c(mean(a), sd(a))); print (c(mean(b), sd(b)))
```


```
[1] 6.294067 3.497769
```

```
[1] -0.10533114 0.05123078
```

```
In [7]: par(mfcol=c(1,1))
        plot(a,b,type="l",xlab=expression(alpha),ylab=expression(beta))
```



```
In [8]: plot(challenger$temp,challenger$oring,pch=19,col="red4",xlab="temperatures",ylab="failures",
for (t in seq(1000,Nsim,le=50)) curve(1/(1+exp(-a[t]-b[t]*x)), add=TRUE,col="grey65",lwd=2.5)
curve(1/(1+exp(-mean(a)-mean(b)*x)),add=TRUE,col="red",lwd=2.5)
postal=rep(0,1000);i=1
```



```
In [9]: for (x0 in c(60,50,40,30)){
  print(c(x0, mean(1/(1+exp(-a-b*x0))),
    sd(1/(1+exp(-a-b*x0))))
}
```

```
[1] 60.0000000 0.4944781 0.1431809
[1] 50.0000000 0.7012076 0.1959169
[1] 40.0000000 0.8134438 0.2077082
[1] 30.0000000 0.8673593 0.1983737
```

1.3 (c)

We want a variance of 100 the previous one, thus we multiply the standard deviation σ by 10.

```

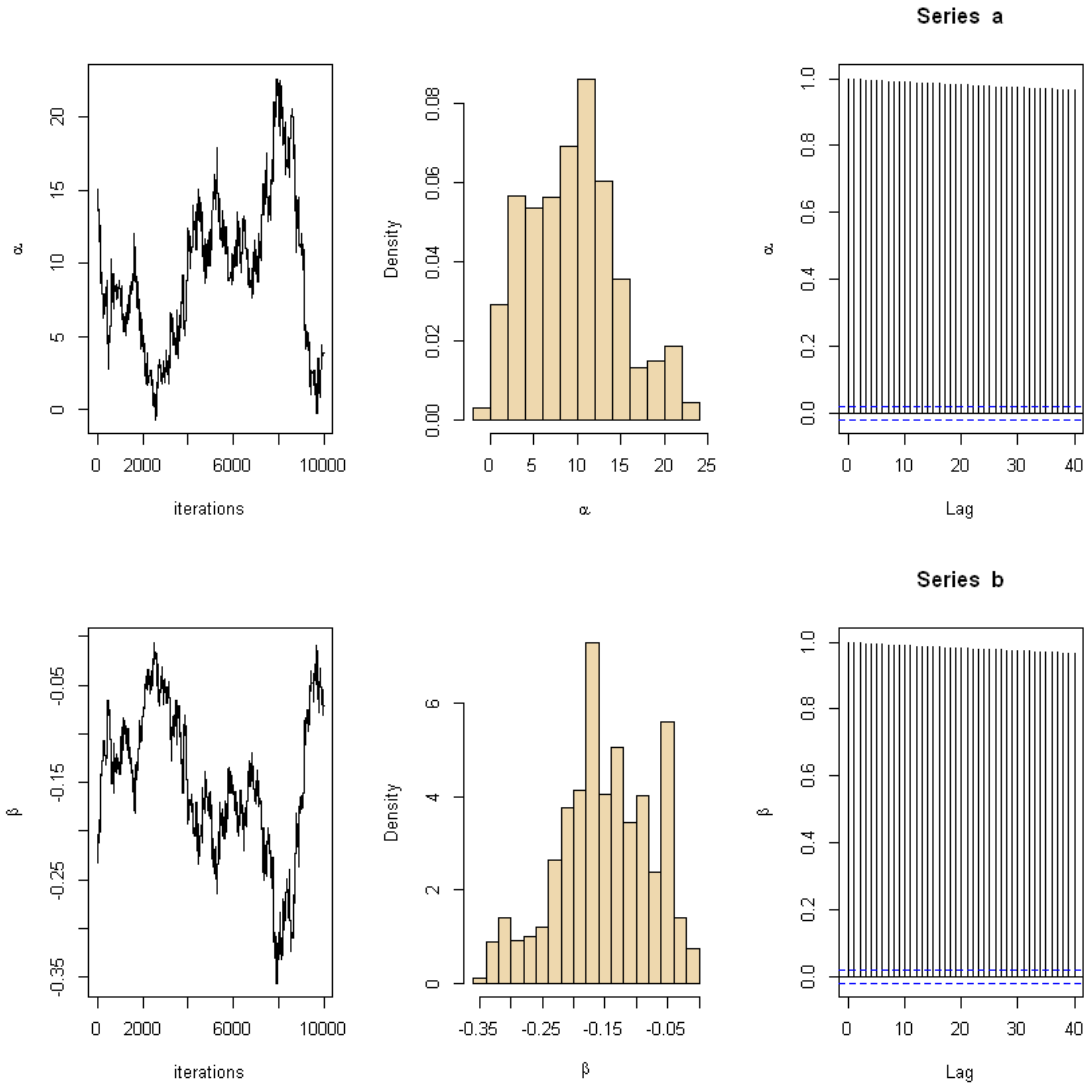
In [10]: set.seed(1)
Nsim=10^4
x=challenger$temp
y=challenger$oring
sigmaa=5*10 ; sigmab=5*10/sd(x)
lpost=function(a,b){sum(y*(a+b*x)-log(1+exp(a+b*x)))+
                        dnorm(a,sd=sigmaa,log=TRUE)+dnorm(b,sd=sigmab,log=TRUE)}
# Initialize a and b to equal the MLEs
beta=as.vector(fit.logistic$coefficients)
a=b=rep(0,Nsim)
a[1]=beta[1]
b[1]=beta[2]
#As scale for the proposal densities consider the square root of the
# cov.unscaled from the ml fit to the logistic model
fit.logistic.summary<-summary(fit.logistic)
scala=sqrt(fit.logistic.summary$cov.unscaled[1,1])
scalb=sqrt(fit.logistic.summary$cov.unscaled[2,2])
for (t in 2:Nsim){
  propa=a[t-1]+sample(c(-1,1),1)*rexp(1)*scala
  if (log(runif(1))<lpost(propa,b[t-1])- lpost(a[t-1],b[t-1]))
    a[t]=propa else
    a[t]=a[t-1]
  propb=b[t-1]+sample(c(-1,1),1)*rexp(1)*scalb
  if (log(runif(1))<lpost(a[t],propb)- lpost(a[t],b[t-1]))
    b[t]=propb
  else b[t]=b[t-1]
}


par(mfrow=c(2,3))
plot(a,type="l",xlab="iterations",ylab=expression(alpha))
hist(a,prob=TRUE,col="wheat2",xlab=expression(alpha),main="")
acf(a,ylab=expression(alpha))
plot(b,type="l",xlab="iterations",ylab=expression(beta))
hist(b,prob=TRUE,col="wheat2",xlab=expression(beta),main="")
acf(b,ylab=expression(beta))


par(mfcol=c(1,1))
plot(a,b,type="l",xlab=expression(alpha),ylab=expression(beta))

plot(challenger$temp,challenger$oring,pch=19,col="red4",xlab="temperatures",ylab="fail")
for (t in seq(1000,Nsim,le=50)) curve(1/(1+exp(-a[t]-b[t]*x)), add=TRUE,col="grey65",lwd=2.5)
curve(1/(1+exp(-mean(a)-mean(b)*x)),add=TRUE,col="red",lwd=2.5)
postal=rep(0,1000);i=1

```







```
In [11]: print(length(unique(a))/Nsim)
          print(length(unique(b))/Nsim)
          print (c(mean(a), sd(a))); print (c(mean(b), sd(b)))
          print(c( mean(1/(1+exp(-a-b*30))),sd(1/(1+exp(-a-b*30)))))
```

```
[1] 0.1062
```

```
[1] 0.1039
```

```
[1] 9.414120 5.119201
```

```
[1] -0.1502712 0.0746917
```

```
[1] 0.9311312 0.1273459
```

From the results above we get:

$\sigma' = \sigma_0$
 $\sigma' = 10\sigma_0$ |---| :---|:---|:---| :---| | α | β | α | β | |acceptance rates | 0.1001|0.0987|
0.1062 |0.1039 |means |6.294067|-0.10533114|9.414120|-0.1502712 |standard deviations
|3.497769|0.05123078|5.119201|0.0746917 |P(error) at 30 F
0.8673593
0.9311312 |standard error
0.1983737
0.1273459

```

In [12]: set.seed(1)
         Nsim=10^4
         x=challenger$temp
         y=challenger$oring
         sigmaa=5*10 ; sigmab=5*10/sd(x)
         lpost=function(a,b){sum(y*(a+b*x)-log(1+exp(a+b*x)))+
                                dnorm(a,sd=sigmaa,log=TRUE)+dnorm(b,sd=sigmab,log=TRUE)}
         # Initialize a and b to equal the MLEs
         beta=as.vector(fit.logistic$coefficients)
         a=b=rep(0,Nsim)
         a[1]=beta[1]
         b[1]=beta[2]
         #As scale for the proposal densities consider the square root of the
         # cov.unscaled from the ml fit to the logistic model
         fit.logistic.summary<-summary(fit.logistic)
         scala=sqrt(fit.logistic.summary$cov.unscaled[1,1])/10
         scalb=sqrt(fit.logistic.summary$cov.unscaled[2,2])/10
         for (t in 2:Nsim){
             propa=a[t-1]+sample(c(-1,1),1)*rexp(1)*scala
             if (log(runif(1))<lpost(propa,b[t-1])- lpost(a[t-1],b[t-1]))
                 a[t]=propa else
                 a[t]=a[t-1]
             propb=b[t-1]+sample(c(-1,1),1)*rexp(1)*scalb
             if (log(runif(1))<lpost(a[t],propb)- lpost(a[t],b[t-1]))
                 b[t]=propb
             else b[t]=b[t-1]
         }

         par(mfrow=c(2,3))
         plot(a,type="l",xlab="iterations",ylab=expression(alpha))
         hist(a,prob=TRUE,col="wheat2",xlab=expression(alpha),main="")
         acf(a,ylab=expression(alpha))
         plot(b,type="l",xlab="iterations",ylab=expression(beta))
         hist(b,prob=TRUE,col="wheat2",xlab=expression(beta),main="")
         acf(b,ylab=expression(beta))

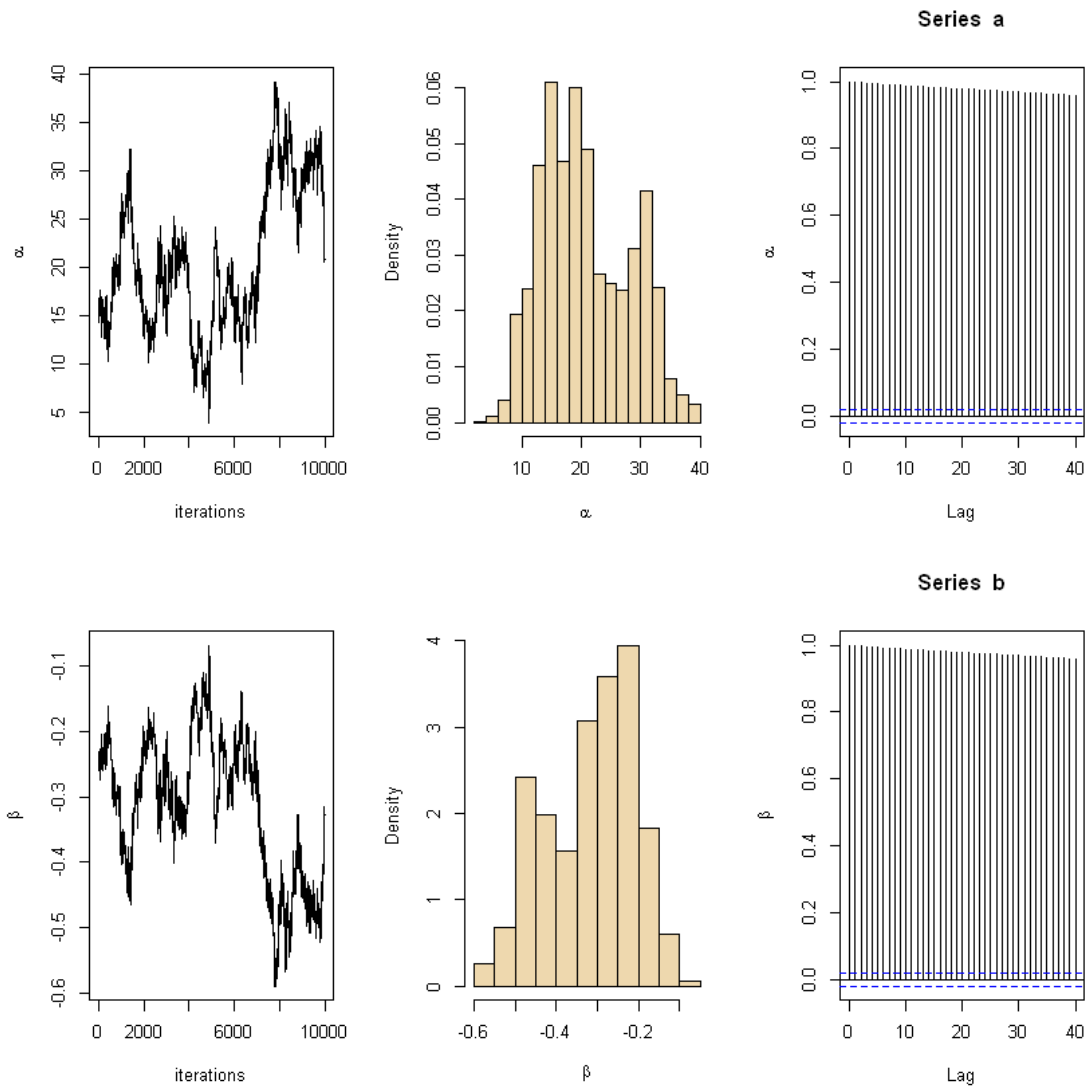
```


```


par(mfcol=c(1,1))
plot(a,b,type="l",xlab=expression(alpha),ylab=expression(beta))

plot(challenger$temp,challenger$oring,pch=19,col="red4",xlab="temperatures",ylab="failures",
for (t in seq(1000,Nsim,le=50)) curve(1/(1+exp(-a[t]-b[t]*x)), add=TRUE,col="grey65"),
curve(1/(1+exp(-mean(a)-mean(b)*x)),add=TRUE,col="red",lwd=2.5)
postal=rep(0,1000);i=1

```







```
In [13]: print(length(unique(a))/Nsim)
          print(length(unique(b))/Nsim)
          print (c(mean(a), sd(a))); print (c(mean(b), sd(b)))
          print(c(mean(1/(1+exp(-a-b*30))),sd(1/(1+exp(-a-b*30)))))
```

```
[1] 0.6104
[1] 0.6141
[1] 20.807518  7.358899
[1] -0.3177214  0.1082698
[1] 0.998901915 0.005897289
```

From the results above we get:

```

|
one-tenth scales
original |---| :---:|:---:|:---: | :---: || $\alpha$ |  $\beta$ | $\alpha$ |  $\beta$ | |acceptance rates | 0.6104|0.6141|
0.1062 |0.1039 |means |20.807518|-0.3177214|9.414120|-0.1502712 |standard deviations
|7.358899|0.1082698|5.119201|0.0746917 |P(error) at 30 F
0.998901915
0.9311312 |standard error
0.005897289
0.1273459
# Problem 2

```

1.4 (a)

The joint posterior distribution simplifies to:

$$\left[\prod_{i=1}^{10} (\lambda_i t_i)^{x_i} e^{-\lambda_i t_i} \lambda_i^{\alpha-1} e^{-\beta \lambda_i} \right] \beta^{10\alpha} \beta^{\gamma-1} e^{-\delta \beta} = \quad (5)$$

$$\left[\prod_{i=1}^{10} \lambda_i^{x_i+\alpha-1} t_i^{x_i} e^{-\lambda_i(t_i+\beta)} \right] \beta^{10\alpha+\gamma-\delta\beta-1} = \quad (6)$$

$$\left[\prod_{i=1}^{10} \lambda_i^{x_i+\alpha-1} t_i^{x_i} \right] e^{-\delta\beta-\sum_{i=1}^{10} \lambda_i(t_i+\beta)} \beta^{10\alpha+\gamma-1} \quad (7)$$

1.5 (b)

The conditional distribution of λ is defined as:

$$\pi(\lambda_i | \beta, t_i, x_i) = \frac{\pi(\lambda_i, \beta | t_i, x_i)}{\pi(\beta)} \quad (8)$$

$$= \frac{\lambda_i^{x_i+\alpha-1} t_i^{x_i} e^{-\delta\beta-\lambda_i(t_i+\beta)} \beta^{10\alpha+\gamma-1}}{\beta^{10\alpha+\gamma-1} e^{-\delta\beta}} = \Gamma(x_i + \alpha, t_i + \beta) \quad (9)$$

The conditional distribution of β will be defined as:

$$\pi(\beta | \lambda_1, \dots, \lambda_{10}) = \Gamma(\gamma + 10\alpha, \delta + \sum_{i=1}^{10} \lambda_i) \quad (10)$$

1.6 (c)


```

In [14]: xdata=c(5, 1, 5, 14, 3, 19, 1, 1, 4, 22)
Time=c(94.32, 15.72, 62.88, 125.76, 5.24, 31.44, 1.05, 1.05, 2.10, 10.48)
nx=length(xdata)
nsim=10^4;alpha = 1.8;gamma= 0.01;delta=1

beta=rgamma(1, shape= alpha, rate = delta)
l1=rgamma(nsim,shape=xdata[1]+alpha,rate=Time[1]+beta)

```

```
In [15]: hist(l1,breaks=50,col="grey",xlab="",main=expression(lambda[1]), freq = F)
```



```
In [16]: summary(l1)
mean(l1)
median(l1)
(var(l1))^0.5
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.008929	0.051490	0.067961	0.071503	0.087785	0.247411

```
0.0715025667320052
0.0679608729590416
```


0.0276862333499326

The maximum likelihood estimates of $\lambda_1, \dots, \lambda_{10}$ will be derived from the MLE for the parameter of the poisson random variable X_i . For a Poisson random variable $Y_i \sim \text{Pois}(\lambda^*)$ the MLE estimate is found as following.

$$\hat{\lambda}_{MLE}^* = \frac{\sum_{i=1}^n Y_i}{n}$$

Thus, for the data we have, we get:


$$\hat{\lambda}_i t_{i_{MLE}} = X_i \hat{\lambda}_{i_{MLE}} = \frac{X_i}{t_i}$$

```
In [17]: xdata=c(5, 1, 5, 14, 3, 19, 1, 1, 4, 22)
         Time=c(94.32, 15.72, 62.88, 125.76, 5.24, 31.44, 1.05, 1.05, 2.10, 10.48)

         l_mle = xdata/Time
         l_mle
```

```
1. 0.053011026293469 2. 0.0636132315521628 3. 0.0795165394402036 4. 0.111323155216285
5. 0.572519083969466 6. 0.604325699745547 7. 0.952380952380952 8. 0.952380952380952
9. 1.9047619047619 10. 2.09923664122137
```

```
In [18]: hist(freq = FALSE, l1,breaks=50,col="grey",xlab="",main=expression(lambda[1]))
         abline(v=l_mle[1], lw=4, col = 'blue')
```



2 (d)

```
In [19]: nx=length(xdata)
nsim=10^4;alpha = 1.8;gamma= 0.01;delta=1
lambda=array(xdata*Time/sum(Time),dim=c(nsim,nx))
beta=rep(gamma*delta,nsim)
for (i in 2:nsim){
  for (j in 1:nx){
    lambda[i,j]=rgamma(1,shape=xdata[j]+alpha,rate=Time[j]+beta[i-1])
  }
  beta[i]=rgamma(1,shape=gamma+nx*alpha,rate=delta+sum(lambda[i,]))
}
```

3 (e)

The 95% credible intervals for the parameters will be calculated as $I_\lambda = [\hat{\lambda}_{MLE} - 1.96\sigma, \hat{\lambda}_{MLE} + 1.96\sigma]$. Thus, for all of the parameters, we have:

```
In [43]: l_int = matrix(nrow = 10, ncol = 2)
  for (i in 1:length(l_mle)){
    sigma = sqrt(var(lambda[,i]))
    l_int[i,1] = l_mle[i]-1.96*sigma
    l_int[i,2] = l_mle[i]+1.96*sigma
  }
```

```
In [45]: l_int

-0.0055884076  0.1116105
-0.1181394103  0.2453659
-0.0031878778  0.1622210
0.0462782423   0.1763681
-0.0002468046  1.1452850
0.3382452347   0.8704062
-0.0732285818  1.9779905
-0.0823859472  1.9871479
0.7829220751   3.0266017
1.3384232977   2.8600500
```

It seems that pump 10 is one of the least reliable ones and pumps 1 and 2 are very reliable.

3.1 (f)

Trying with a prior with parameters computed from the MLE estimates. For gamma distributions $X \sim \Gamma(k, \theta)$, we have:

$$\hat{\theta}_{MLE} = \frac{1}{kN} \sum_{i=1}^N x_i s = \ln\left(\frac{1}{N} \sum_{i=1}^N x_i\right) - \frac{1}{N} \sum_{i=1}^N \ln(x_i) k \simeq \frac{3 - s + \sqrt{(s-3)^2 + 24s}}{12s}$$

```
In [62]: xdata=c(5, 1, 5, 14, 3, 19, 1, 1, 4, 22)
  Time=c(94.32, 15.72, 62.88, 125.76, 5.24, 31.44, 1.05, 1.05, 2.10, 10.48)
  nx=length(xdata)
  nsim=10^4;alpha = 1.8;gamma= 0.01;delta=1


  s = log(sum(xdata)/nx)-sum(log(xdata))/nx
  alpha = (3-s+sqrt((s-3)^2+24*s))/(12*s)

  beta=rgamma(1, shape= alpha, rate = delta)
  l1=rgamma(nsim,shape=xdata[1]+alpha,rate=Time[1]+beta)
  alpha

0.99116678629947
```

```
In [58]: hist(l1,breaks=50,col="grey",xlab="",main=expression(lambda[1]), freq = F)
```

λ_1



```
In [59]: summary(l1)
         mean(l1)
         median(l1)
         (var(l1))^0.5
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.008391	0.044136	0.059198	0.062665	0.077269	0.217807

```
0.0626651229839855
0.0591984331256458
0.0256608479223233
```

The maximum likelihood estimates of $\lambda_1, \dots, \lambda_{10}$ will be derived from the MLE for the parameter of the poisson random variable X_i . For a Poisson random variable $Y_i \sim \text{Pois}(\lambda^*)$ the MLE estimate is found as following.

$$\hat{\lambda}_{MLE}^* = \frac{\sum_{i=1}^n Y_i}{n}$$

Thus, for the data we have, we get:


$$\hat{\lambda}_i t_{i_{MLE}} = X_i \hat{\lambda}_{i_{MLE}} = \frac{X_i}{t_i}$$

```
In [60]: xdata=c(5, 1, 5, 14, 3, 19, 1, 1, 4, 22)
         Time=c(94.32, 15.72, 62.88, 125.76, 5.24, 31.44, 1.05, 1.05, 2.10, 10.48)

         l_mle = xdata/Time
         l_mle

1. 0.053011026293469 2. 0.0636132315521628 3. 0.0795165394402036 4. 0.111323155216285
5. 0.572519083969466 6. 0.604325699745547 7. 0.952380952380952 8. 0.952380952380952
9. 1.9047619047619 10. 2.09923664122137

In [61]: hist(freq = FALSE, l1,breaks=50,col="grey",xlab="",main=expression(lambda[1]))
         abline(v=l_mle[1], lw=4, col = 'blue')
```



4 (d)

```
In [63]: nx=length(xdata)
nsim=10^4;alpha = 1.8;gamma= 0.01;delta=1
lambda=array(xdata*Time/sum(Time),dim=c(nsim,nx))
beta=rep(gamma*delta,nsim)
for (i in 2:nsim){
  for (j in 1:nx){
    lambda[i,j]=rgamma(1,shape=xdata[j]+alpha,rate=Time[j]+beta[i-1])
  }
  beta[i]=rgamma(1,shape=gamma+nx*alpha,rate=delta+sum(lambda[i,]))
}
```

The 95% credible intervals for the parameters will be calculated as $I_\lambda = [\hat{\lambda}_{MLE} - 1.96\sigma, \hat{\lambda}_{MLE} + 1.96\sigma]$. Thus, for all of the parameters, we have:

```
In [64]: l_int = matrix(nrow = 10, ncol = 2)
         for (i in 1:length(l_mle)){
           sigma = sqrt(var(lambda[,i]))
           l_int[i,1] = l_mle[i]-1.96*sigma
           l_int[i,2] = l_mle[i]+1.96*sigma
         }
```

```
In [45]: l_int
```

```
-0.0055884076  0.1116105
-0.1181394103  0.2453659
-0.0031878778  0.1622210
 0.0462782423  0.1763681
-0.0002468046  1.1452850
 0.3382452347  0.8704062
-0.0732285818  1.9779905
-0.0823859472  1.9871479
 0.7829220751  3.0266017
 1.3384232977  2.8600500
```